

Parallel Sorting Algorithms

(cf. Grama et al.)

Sorting: Overview

- Definitely one of the most important and used topics in computing
- Sorting can be **comparison-based** or **non-comparison-based**
- The fundamental operation of comparison-based sorting is the **exchange-comparison** (compare-exchange)
- The **lower limit** of any sequential sorting algorithm for n numbers is $\Theta(n \log n)$
- In the parallel computing field, the **most** we can expect with n processors is $O(\log n)$! In fact it is obtained, but with very large constants

Example: Sorting of n numbers

Input: sequence of **n** numbers $\langle a_1, a_2, a_3, \dots, a_n \rangle$

Bubble Sort

Output: permutation $\langle a'_1, a'_2, a'_3, \dots, a'_n \rangle$ of the elements so that $a'_1 \leq a'_2, \leq a'_3 \leq \dots \leq a'_n$

```
/* Bubble sort for integers */
#define SWAP(a,b) { int t; t=a; a=b; b=t; }
void SORT( int a[], int n )
/* Pre-condition: a contains n items to be sorted */
{
    int i, j;
    /* Make n passes through the array */
    for(i=0; i<n; i++)
    {
        /* From the first element to the end of the unsorted section */
        for(j=1; j<(n-i); j++)
        {
            /* If adjacent items are out of order, swap them */
            if( a[j-1]>a[j] ) SWAP(a[j-1], a[j]);
        }
    }
}
```

Example: Parallel Sorting

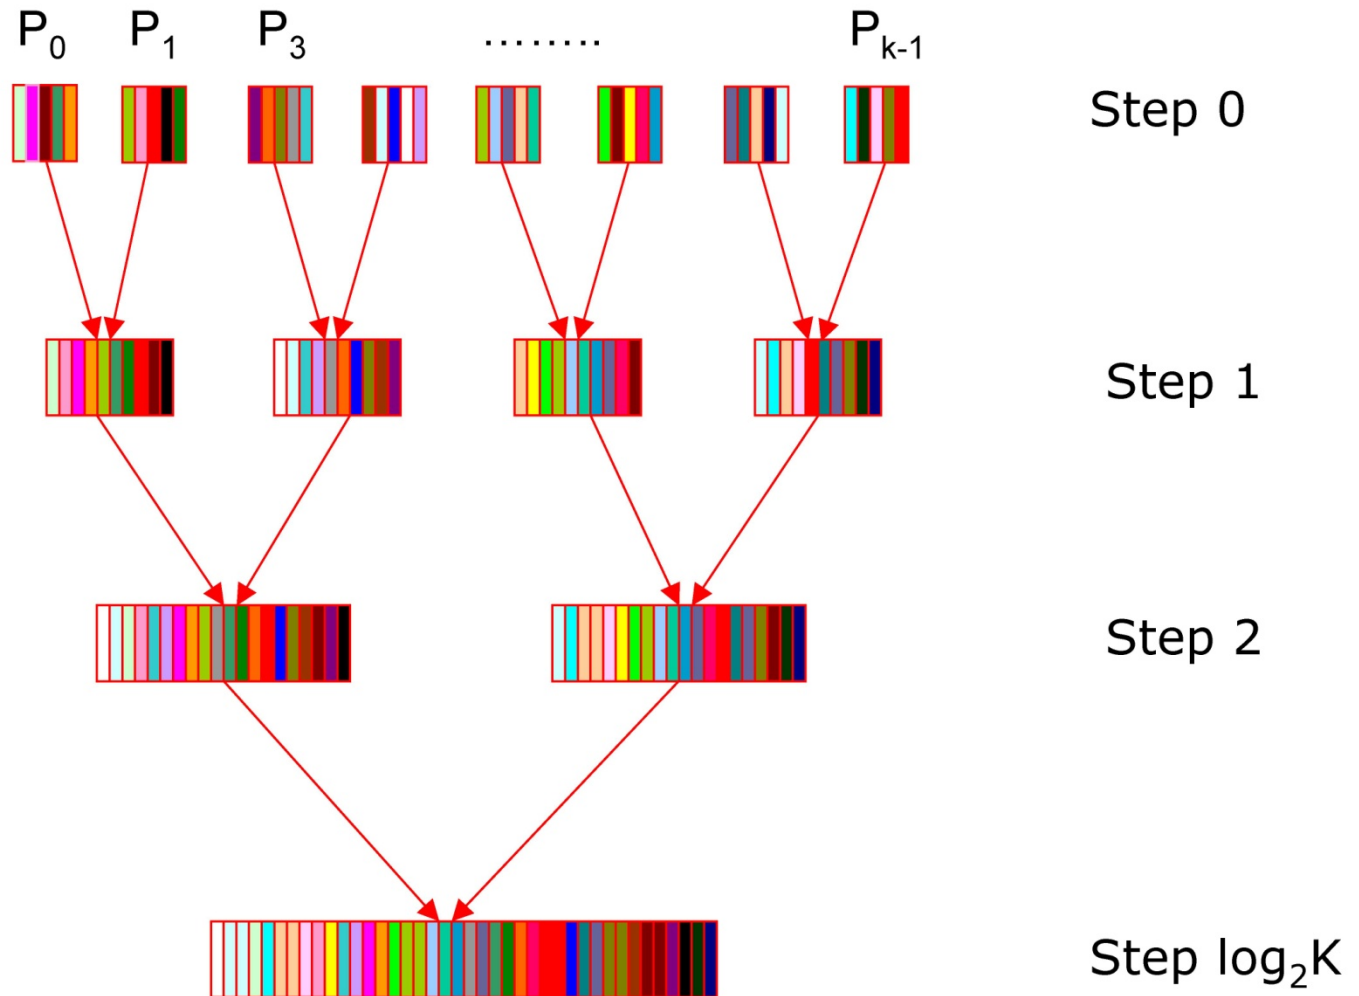
Idea !

```
SORT(a[0 : n/2-1])
```

```
SORT(a[n/2 : n-1])
```

```
MERGE(a[0 : n/2-1], a[n/2 : n-1])
```

Example: Parallel Sorting



Formally

- Let $S = \langle a_1, a_2, \dots, a_n \rangle$ be a sequence n elements
- **The ordering of S** consists in a sequence $S' = \langle a'_1, a'_2, \dots, a'_n \rangle$ such that $a'_i \leq a'_j$ for $1 \leq i \leq j \leq n$ and S' is a permutation of S

Compare exchange

- Base of each **comparison based** algorithm

```
if (A > B) {  
    temp = A;  
    A = B;  
    B = temp;  
}
```

compare-exchange operation

Sorting: Basics


What is a **ordered parallel sequence**? Where are the input and output lists stored?

The main assumption is that the input and output lists are **distributed**

- By convention, the sorted list is **partitioned** with the property that each **sub-list is sorted** and each element of the list of processor P_i is smaller than those in P_j if $i < j$
- The previous fact is at the basis of every parallel sorting algorithm!

Sorting: Parallel Compare Exchange Operation

In **sequential sorting** algorithms, the input and any ordered sequences (including the output) are stored in the memory of the **same process**. In the parallel case, the data can, of course, be stored in other processes, making the algorithm further complicated!




A compare-exchange operation (one element per processor). The processes P_i and P_j send their elements to each other. The process P_i keeps $\min\{a_i, a_j\}$, and P_j keeps $\max\{a_i, a_j\}$

Sorting: Basics

What is the counterpart of a parallel sequence comparison?

- If each processor has an element, the **compare-exchange** operation stores the smallest element in the processor with **smaller id**. This is done in time $t_s + t_w$
- If you have more than one element per processor, this operation is called **compare split**. Suppose that each of the two processors have n / p elements.
- **After the compare-split operation, the n/p smallest elements are in the processor P_i while the greater n / p are in P_j , where $i < j$**
- The time for a compare-split operation is $(t_s + t_w n/p)$, assuming that the two "partial" lists were initially ordered

Sorting: Compare Split Parallel Operation



A compare-split operation. Each process sends its block of size n / p to another process. Each process "merges" the block that receives with its own and maintains only the appropriate half. In this example, the process P_i maintains the minor elements and the process P_j maintains the bigger.

Sorting Networks

- They are "**comparator**" networks, designed for sorting
- A **comparator is a device** with two inputs x and y and two outputs x' and y' . For an increasing comparator, we have $x' = \min \{x, y\}$ and $y' = \max \{x, y\}$; and vice-versa.
- We denote an increasing comparator with \oplus and decreasing comparator with \ominus
- The "**speed**" of the network is proportional to its depth

Sorting Networks: Comparators







Figure 9.3 A schematic representation of comparators: (a) an increasing comparator, and (b) a decreasing comparator.

Ordering Networks

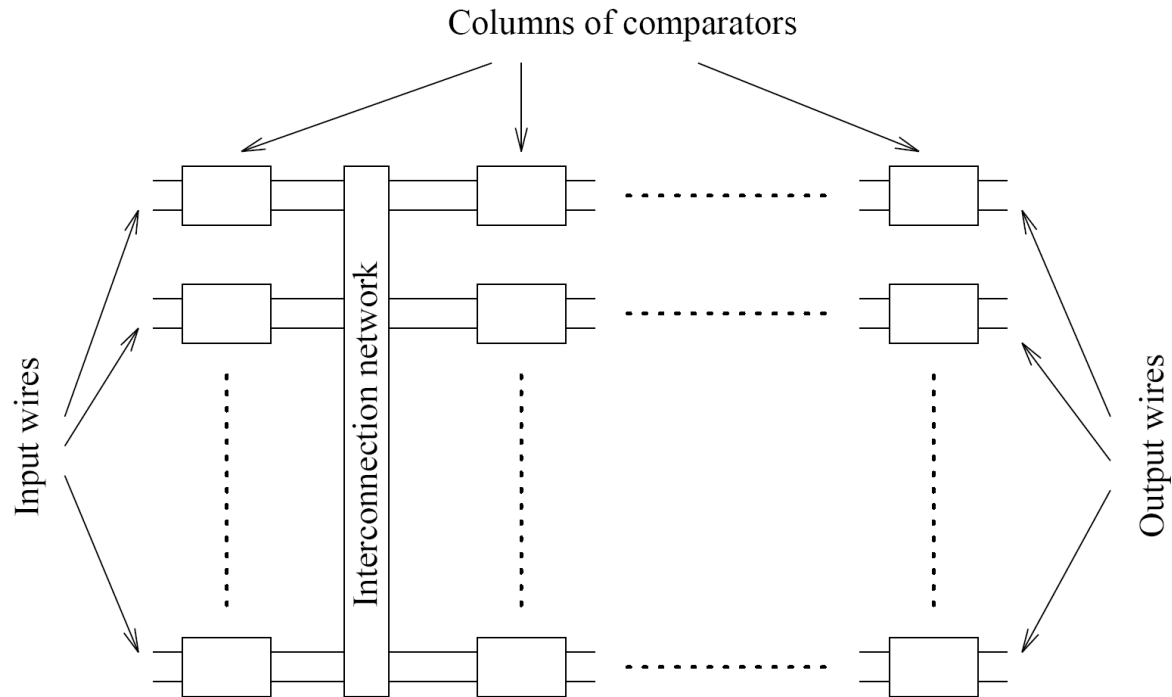


Figure 9.4 A typical sorting network. Every sorting network is made up of a series of columns, and each column contains a number of comparators connected in parallel.

Sorting Networks : Bitonic Sort

- A **bitonic** sorting network sorts n elements in time $\Theta(\log^2 n)$
- A **bitonic sequence** has two “**tonalities**” - ascending and descending, or vice versa. Any cyclic rotation of such networks is also considered **bitonic**
- For example $\langle 1, 2, 4, 7, 6, 0 \rangle$ is a bitonic sequence, because before it increases and then decreases. $\langle 8, 9, 2, 1, 0, 4 \rangle$ is another bitonic sequence, because it is a cyclic shift of $\langle 0, 4, 8, 9, 2, 1 \rangle$
- The heart of the network is the reordering of the bitonic sequence into an ordered sequence

Sorting Networks : Bitonic Sort

- Let $s = \langle a_0, a_1, \dots, a_{n-1} \rangle$ be a bitonic sequence **such that**
 $a_0 \leq a_1 \leq \dots \leq a_{n/2-1}$ **e** $a_{n/2} \geq a_{n/2+1} \geq \dots \geq a_{n-1}$.
- Let's consider the following sub-sequences of s :
$$s_1 = \langle \min\{a_0, a_{n/2}\}, \min\{a_1, a_{n/2+1}\}, \dots, \min\{a_{n/2-1}, a_{n-1}\} \rangle$$
$$s_2 = \langle \max\{a_0, a_{n/2}\}, \max\{a_1, a_{n/2+1}\}, \dots, \max\{a_{n/2-1}, a_{n-1}\} \rangle$$
- Note that s_1 and s_2 are both bitonic and that each element of s_1 is smaller than each element of s_2 .
- We can re-apply the procedure recursively on s_1 and s_2 to obtain subsequences of size 1, i.e. the **ordered sequence (bitonic merge)**

Sorting Networks : Bitonic Sort

Original																
sequence	3	5	8	9	10	12	14	20	95	90	60	40	35	23	18	0
1st Split	3	5	8	9	10	12	14	0	95	90	60	40	35	23	18	20
2nd Split	3	5	8	0	10	12	14	9	35	23	18	20	95	90	60	40
3rd Split	3	0	8	5	10	9	14	12	18	20	35	23	60	40	95	90
4th Split	0	3	5	8	9	10	12	14	18	20	23	35	40	60	90	95

Figure 9.5 Merging a 16-element bitonic sequence through a series of $\log 16$ bitonic splits.

The sorting of a bitonic sequence occurs in $\log n$ steps!

Sorting Networks : Bitonic Sort

- We can now build a **sorting** network to implement the **bitonic merge algorithm**.
- This network is called ***bitonic merging network***.
- This network contains **$\log n$ columns**. Each column contains $n / 2$ **comparators** and performs one step of the bitonic merge.
- We denote by $\oplus\text{BM}[n]$ a **bitonic merging network** with n inputs
- **Replacing** the comparators \oplus with the \ominus comparators we get a **descending output sequence** : This network is called $\ominus\text{BM}[n]$
- $\oplus\text{BM}[n]$ is also known as **Co-Ex-Low**; $\ominus\text{BM}[n]$ is also known as **Co-Ex-High**

Sorting Networks : Bitonic Sort

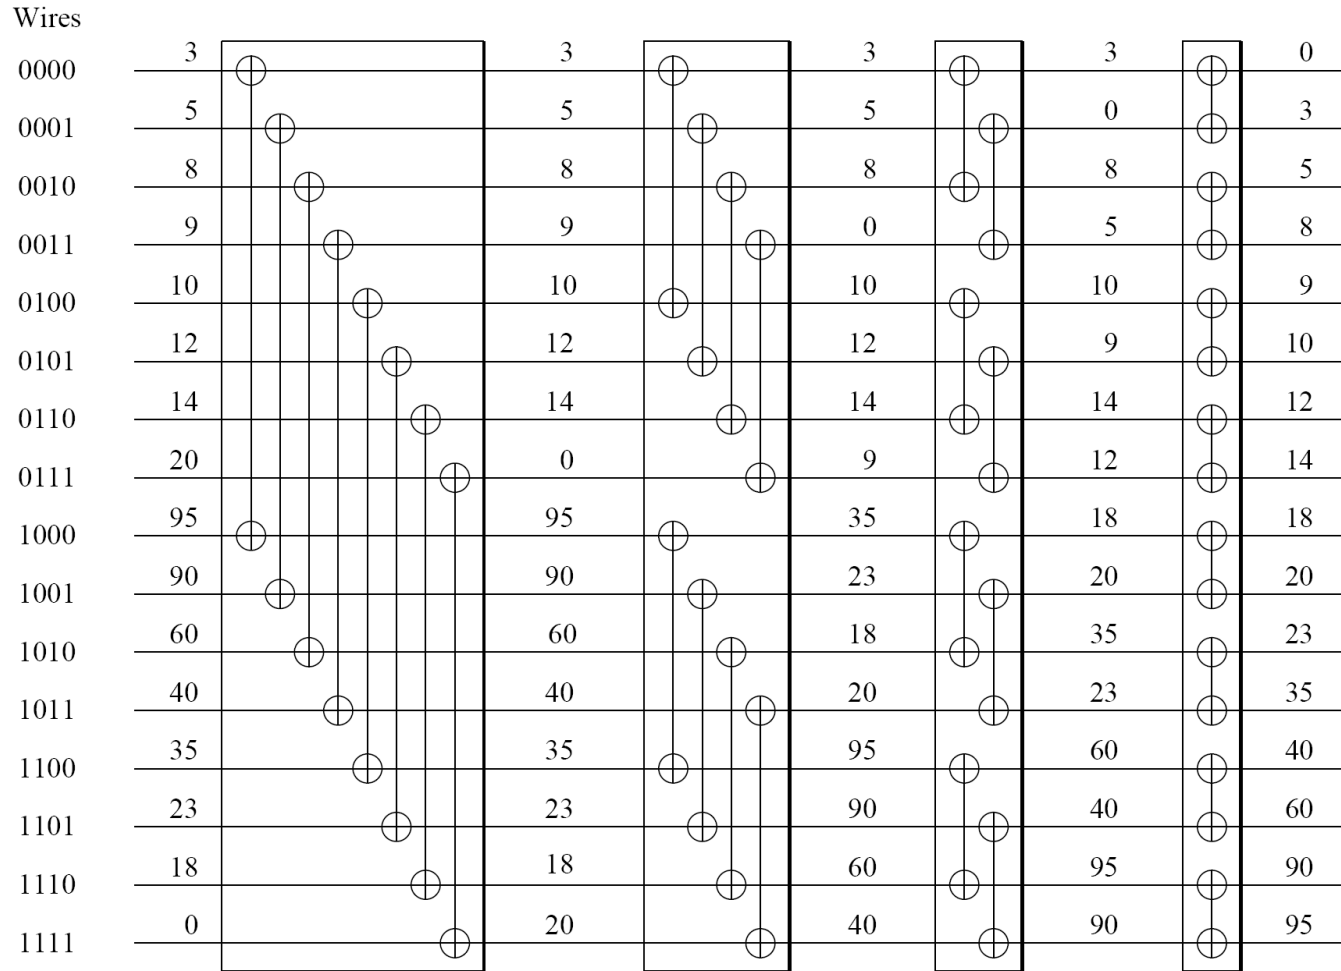


Figure 9.6 A bitonic merging network for $n = 16$. The input wires are numbered $0, 1, \dots, n - 1$, and the binary representation of these numbers is shown. Each column of comparators is drawn separately; the entire figure represents a $\oplus\text{BM}[16]$ bitonic merging network. The network takes a bitonic sequence and outputs it in sorted order.

Sorting Networks : Bitonic Sort

*How do we sort a not ordered sequence
using a bitonic merge?*

- We must first build a single bitonic sequence from the given sequence
- **We note that a sequence of length 2 is bitonic.** So any unordered sequence of elements can be thought of as a **concatenation** of bitonic sequences!
- A bitonic sequence of length 4 can be constructed by ordering the first two elements using $\oplus\text{BM}[2]$ and the next two with $\ominus\text{BM}[2]$
- This process can be **repeated** to generate more bitonic sequences

Sorting Networks : Bitonic Sort

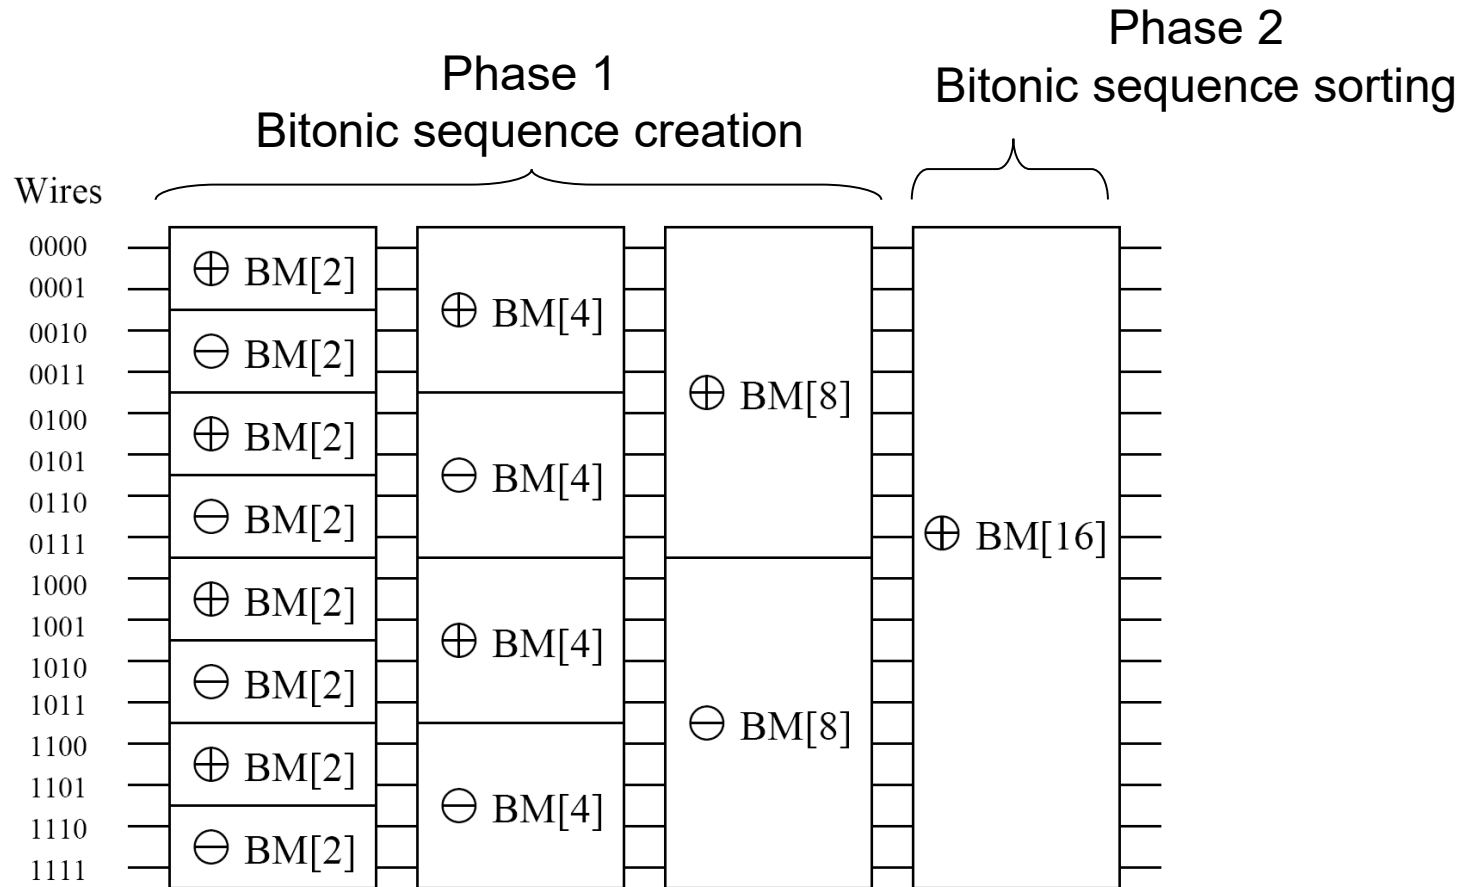


Figure 9.7 A schematic representation of a network that converts an input sequence into a bitonic sequence. In this example, $\oplus\text{BM}[k]$ and $\ominus\text{BM}[k]$ denote bitonic merging networks of input size k that use \oplus and \ominus comparators, respectively. The last merging network ($\oplus\text{BM}[16]$) sorts the input. In this example, $n = 16$.

Sorting Networks : Bitonic Sort (Phase 1)

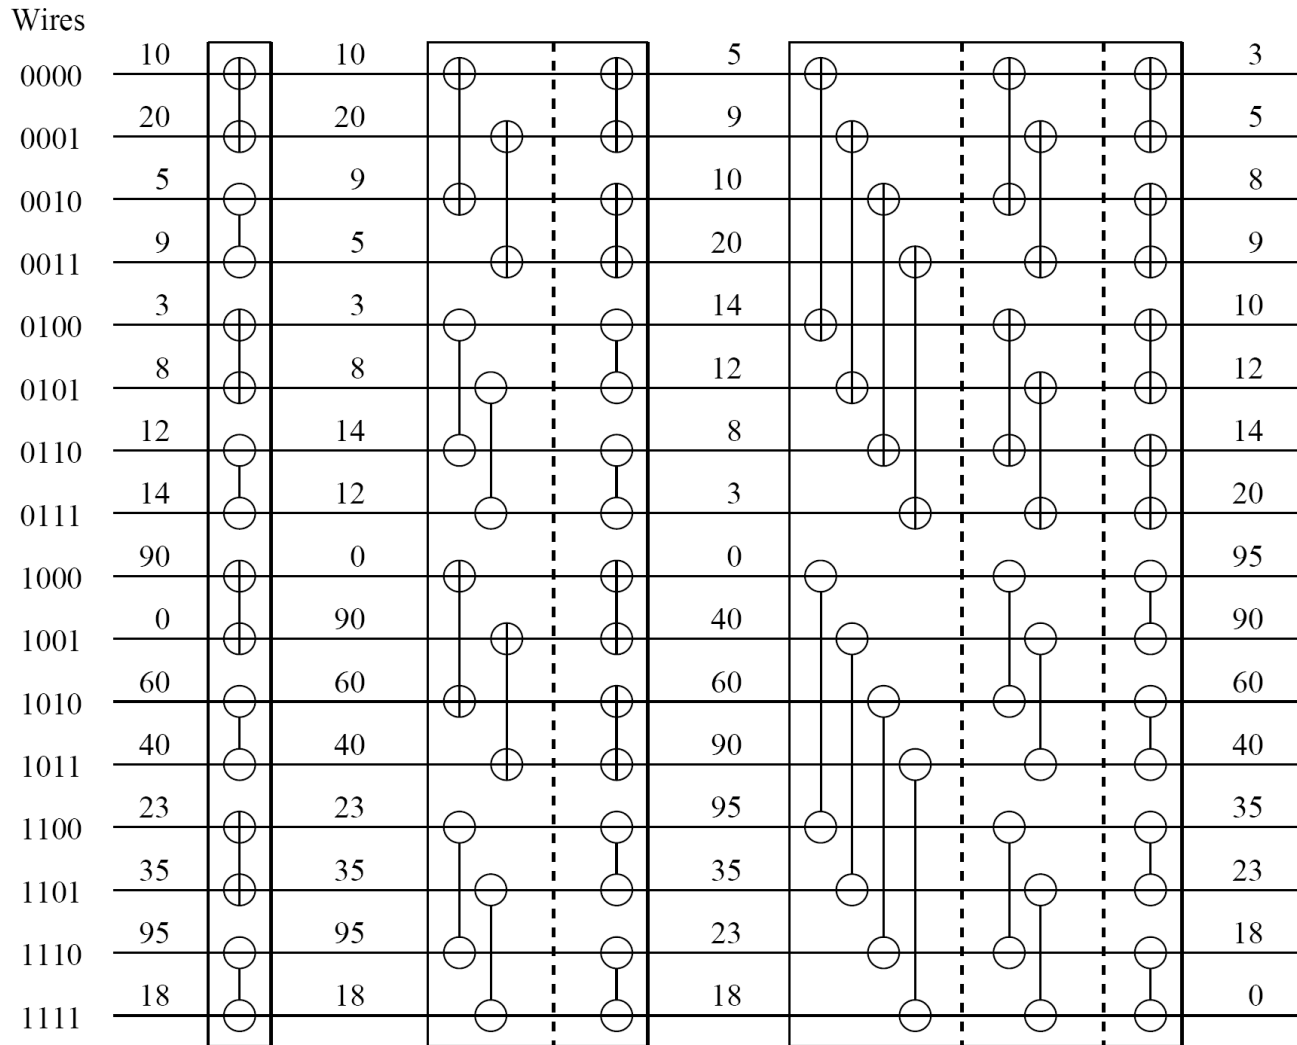


Figure 9.8 The comparator network that transforms an input sequence of 16 unordered numbers into a bitonic sequence. In contrast to Figure 9.6, the columns of comparators in each bitonic merging network are drawn in a single box, separated by a dashed line.

Sorting Networks : Bitonic Sort

- So, the **complete algorithm (GENERAL BITONIC SORT)** consists in:
 - **Creation of a bitonic sequence** starting from a non order one (fig. 9.8) (Phase 1)
 - Application of the **bitonic merging network** to such sequence (fig. 9.6), with depth $\log n$ (Phase 2 – Bitonic sort)
- Phase 1 totally orders $n/2$ elements
- So, resolving the following recurrence relation (d = tree depth)

$$d(n) = d\left(\frac{n}{2}\right) + \log n$$

$$d(n) = \sum_{i=1}^{\log n} i = \frac{\log^2 n + \log n}{2} = \Theta(\log^2 n)$$

Sorting Networks : Bitonic Sort

$$\begin{aligned}T(n) &= \log(n) + \log(n/2) + \log(n/4) + \dots + 1 = \\&= \log(n) + [\log(n)-1] + [\log(n)-2] + [\log(n) - 3] + \dots + 1 = \\&= \log(n)*\log(n) - (2 + 3 + \dots \log(n)) = \\&= \log(n)*\log(n) - ((2+\log(n))*\log(n))/2 = \\&= \log(n)*\log(n) - \sim 0.5\log(n)*\log(n) = \\&= \sim 0.5\log(n)*\log(n)\end{aligned}$$


So,

$$\mathbf{T(n) = O(\log(n)*\log(n))}$$


Bitonic Sort for hypercubes

- Consider the case in which we have **one element** for processor
- Note that the compare-exchange operation takes place between two "connections" (therefore, processors) that differ by exactly one bit!
- It's therefore **natural** to implement a connections-processors direct mapping using a hypercube where, let's recall, two processes are "close" if their binary representations differ by exactly one bit!


Bitonic Sort for hypercubes




Step 1



Step 2




Step 3



Step 4

Communications during the **last step** of the bitonic sort.
Each connection is mapped to a process of the hypercube; each connection is a compare-exchange between processes

Bitonic Sort for hypercubes



Features of the communications in bitonic sort on a hypercube. During each phase of the algorithm, processors communicate in dimensions shown in figure

Bitonic Sort for hypercubes

```
1.  procedure BITONIC_SORT(label, d)
2.  begin
3.      for i := 0 to d - 1 do
4.          for j := i downto 0 do
5.              if (i + 1)st bit of label ≠ jth bit of label then
6.                  comp_exchange_max(j);
7.              else
8.                  comp_exchange_min(j);
9.  end BITONIC_SORT
```

Parallel formulation of bitonic sort on a hypercube with $n = 2^d$ processes

Bitonic Sort for hypercubes

- During each phase of the algorithm, each process executes a compare-exchange operation (single communication to the “direct” neighbour)
- Since each step takes $\Theta(1)$, the parallel time is:


$$T_p = \Theta(\log^2 n)$$

- The algorithm is cost optimal with respect to the sequential version, but not with respect to the best parallel sorting algorithm


Bitonic Sort for Grids

- The connectivity of a grid / mesh is lower than that of a hypercube, so we expect some overhead ...
- A shuffled row-major mapping is usually considered on processes ...


Mapping Bitonic Sort to Meshes



(a)




(b)



(c)

Different ways of mapping the input wires of the bitonic sorting network to a mesh of processes: (a) row-major mapping, (b) row-major snakelike mapping, and (c) row-major shuffled mapping

Mapping Bitonic Sort to Meshes



The **last stage** of the bitonic sort algorithm for $n = 16$ on a mesh, using the row-major shuffled mapping. During each step, process pairs compare-exchange their elements. Arrows indicate the pairs of processes that perform compare-exchange operations.

Block of Elements Per Processor


- Each process is assigned a block of n/p elements.
- The first step is a local sort of the local block (v. applicazioni sequenti)
- **Each subsequent compare-exchange operation is replaced by a compare-split operation**
- We can effectively view the bitonic network as having $(1 + \log p)(\log p)/2$ steps.

Applications/ Appendix


General Bitonic Sort: Phase 1 (Example)

To transform a vector of 8 elements in a bitonic one, partition it in 4 pairs, or in 4 bitonic sub-sequences of length = 2 and order them so you have 2 bitonic subsequences of length = 4

Initial vector:



4 subsequences of length 2



General Bitonic Sort: Phase 1 (Example)

Co-Ex-Lo

6 3 | 9 8 | 1 5 | 4 2

Co-Ex-Hi

3 6 | 9 8 | 1 5 | 4 2


We obtain:

3 6 | 9 8 | 1 5 | 4 2

Bitonic sequence
 $l=4$

Bitonic sequence
 $l=4$

General Bitonic Sort: Phase 1 (Example)




We order each of the two subvectors (with the bitonic Sort -1 phase)
in order to have a single bitonic vector




The first in ascending order
the second in descending order

General Bitonic Sort: Phase 1 (Example)

We have:





Ordered
subsequence in
increasing order




Ordered
subsequence in
decreasing order

General Bitonic Sort: Phase 2 (Example)

- **Last step:** Sorting a bitonic sequence of length 8 obtained before (with bitonic SORT):



After $\log n$ steps...



Global ordered
list

Sorting in Parallel: General Example

- How can we implement a parallel sorting algorithm, starting from any **sequential** algorithm?

Example $p=2$

Let's distribute the vector

➤ $L = (25, 7, 1, 9, 81, 3, 28, 12, 6, 20)$

25
7
1
9
81

3
28
12
6
20

General Strategy

- Phase 1 (Local Sort)
 - Each processor sorts its vector with a standard sorting algorithm (e.g., quicksort)
- Phase 2 (Merging)
 - The (sub)vectors are opportunely combined to obtain an ordered global vector

Phase 1 (Local sort)

P0

25
7
1
9
81

P1

3
28
12
6
20

Each process sorts in increasing manner
its vector

P0

1
7
9
25
81

P1

3
6
12
20
28

Last phase, local sort

Phase 2 (Merging)

The compare-exchange operation is applied to the distributed sequence between the 2 processes



Low process

Process P0 contains the
smallest elements

Co-Ex-Lo




High process

Process P1 contains the
biggest elements


Co-Ex-Hi

I PASSO: step 1.0 → COMUNICAZIONE - Compare Split?



Step 1.0: P0 sends to P1 its max= 81, while P1 sends to P0 its min=3

I PASSO : step 1.1 → CONFRONTO




IN PARALLELO

Step 1.1: P0 compares 3 with 81, if smaller it inserts it

P1 compares 81 with 3, if greater it inserts it


I PASSO: step 1.2 → INSERIMENTO



Step 1.2: Each process


- 1) Inserts the received number
- 2) Erases the sent number
- 3) Reorders vector

II PASSO: step 2.0 → COMUNICAZIONE



Step 1.0: P_0 invia a P_1 il proprio max = 25,
mentre P_1 invia a P_0 il proprio min = 6

II PASSO: step 2.1 → CONFRONTO




IN PARALLELO

Step 1.1: P_0 confronta 6 con 25: se è minore lo inserisce;

P_1 confronta 25 con 6: se è maggiore lo inserisce

II PASSO: step 2.2 → INSERIMENTO




IN PARALLELO

Step 1.2: Ciascun processore

- 1) inserisce il numero ricevuto
- 2) Elimina il numero inviato
- 3) riordina il vettore


III PASSO: step 3.0 → COMUNICAZIONE

FASE 2 (MERGING)



Step 1.0: P_0 invia a P_1 il proprio max = 9,
mentre P_1 invia a P_0 il proprio min = 12

III PASSO: step 3.1 → CONFRONTO



Step 1.1: P_0 confronta il numero **12** con 9: non è minore e termina l'esecuzione.

P_1 confronta il numero **9** con 12: non è maggiore e termina l'esecuzione

PGBS: Ordinamento parallelo di un vettore

P0

1
3
6
7
9

P1


12
20
25
28
81

I 5 numeri più **piccoli** sono in P_0
mentre i 5 più **grandi** sono in P_1

Le due sottoliste iniziali sono globalmente ordinate
secondo l'ordine crescente dell'identificativo dei processi

Esempio: $p=4$, $n=20$

Passo 1: distribuzione dati



Esempio: $p=4$, $n=20$

FASE 1 (SORTING LOCALE)

P_0	P_1	P_2	P_3
2	1	6	3
4	5	8	7
19	9	11	17
29	15	18	20
34	23	38	22

Ogni processore ordina in modo **crescente**
la propria **sottolista**

Esempio: $p=4$, $n=20$

FASE 2 (MERGING)

P_0

P_1

P_2

P_3

2
4
19
29
34

1
5
9
15
23


6
8
11
18
38

3
7
17
20
22

Si applica l'algoritmo GBS a tutto il vettore


Esempio: $p=4$, $n=20$

FASE 2 (MERGING)



Esempio: $p=4$, $n=20$

FASE 2 (MERGING)



Fine passo $i=1$ e $j=1$
Il vettore è bitonico

Esempio: $p=4$, $n=20$

FASE 2 (MERGING)


P_0	P_1	\vdots	P_2	P_3
1	15		38	11
2	19		22	8
4	23		20	7
5	29		18	6
9	34		17	3

\vdots
Passo $i=2$

ordiniamo globalmente il vettore

Esempio: $p=4$, $n=20$

FASE 2 (MERGING)




Passo $i=2$ ($j=1$)

Esempio: $p=4$, $n=20$

FASE 2 (MERGE)

La lista è ora globalmente ordinata




Passo $i=2$ ($j=2$)


Esempio: $p=4$, $n=20$

La lista è ora globalmente ordinata


P_0





P_1



P_2



P_3



More Appendix

Bubble Sort

First, largest number moved to the end of list by a series of compares and exchanges, starting at the opposite end.

Actions repeated with subsequent numbers, stopping just before the previously positioned number.

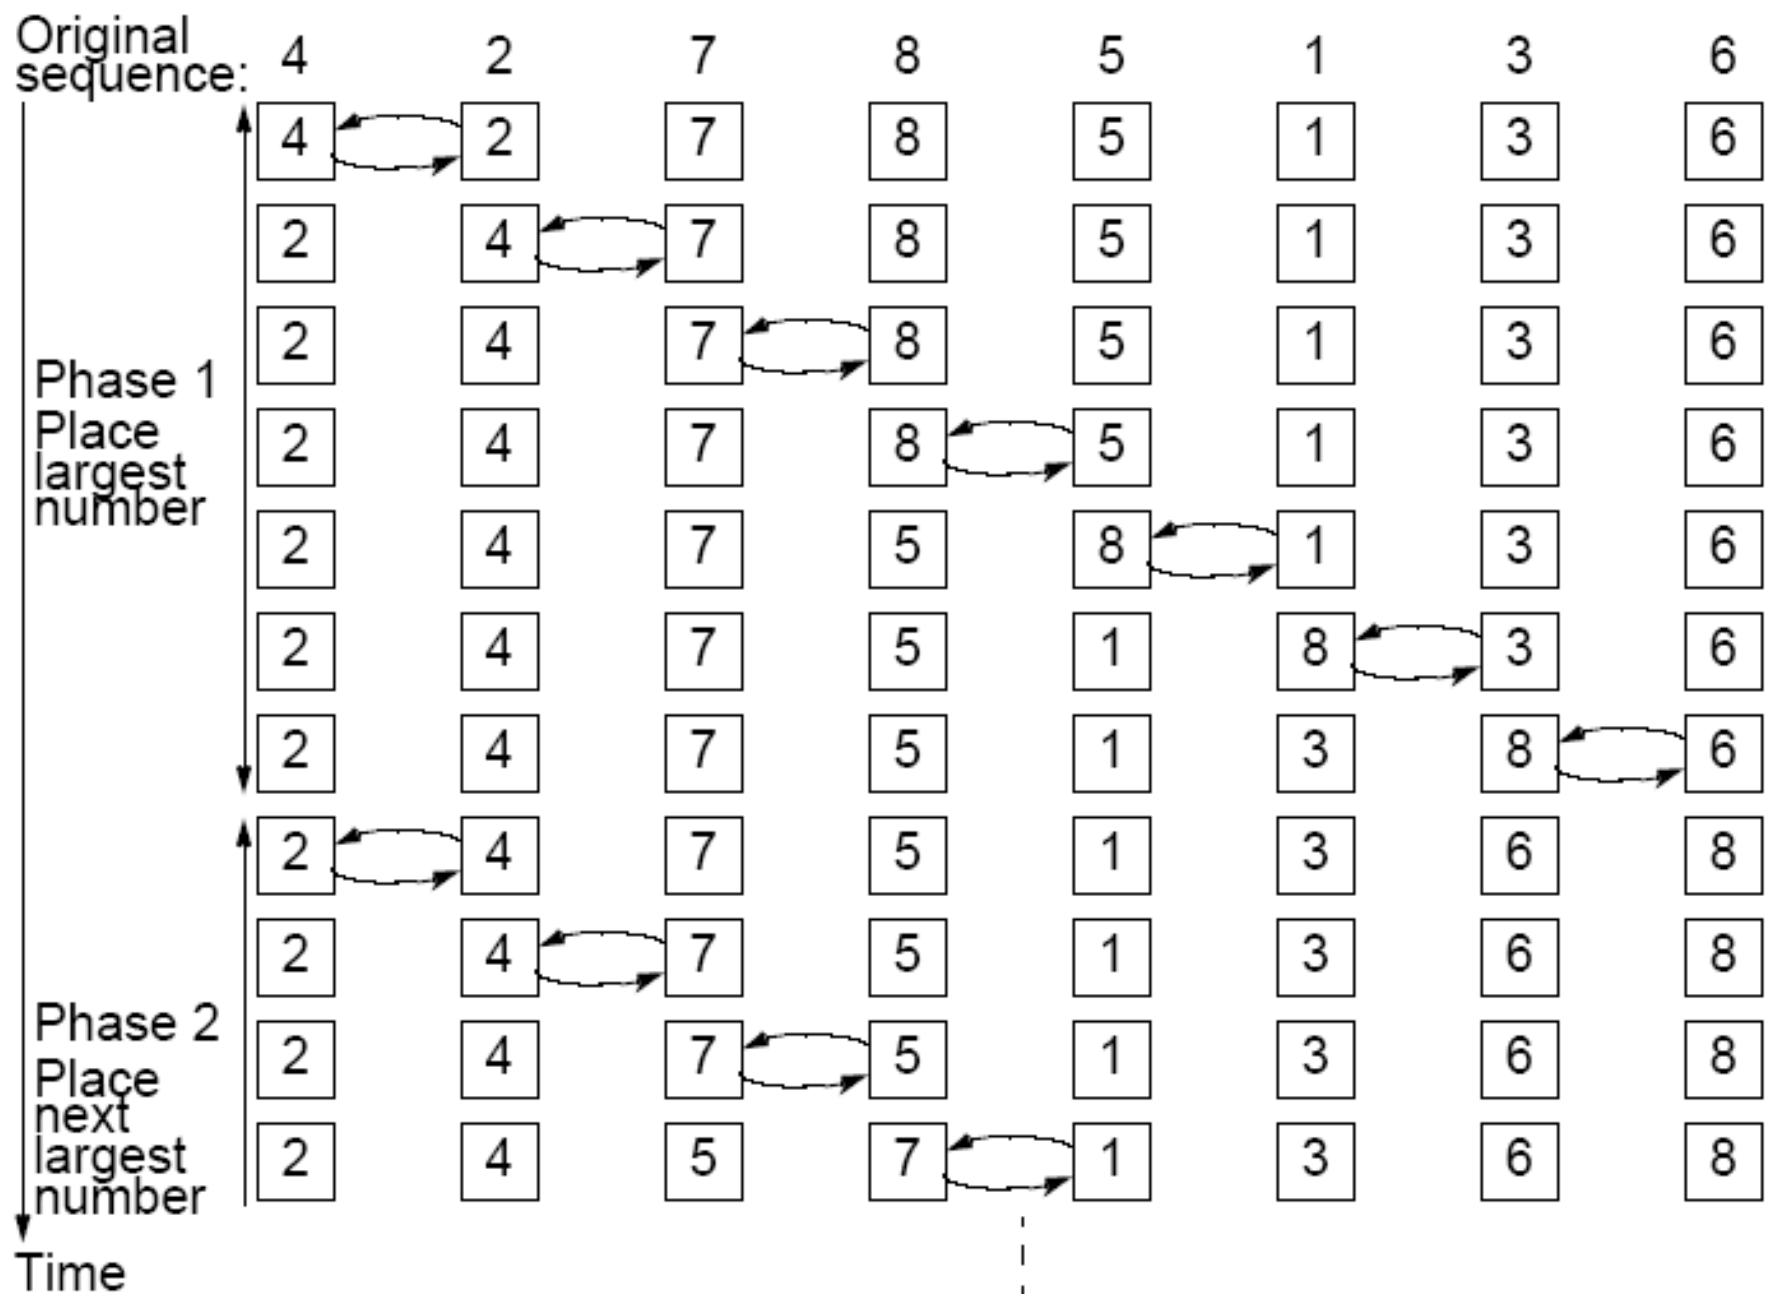
In this way, the larger numbers move (“bubble”) toward one end,

Bubble Sort and its Variants

The sequential bubble sort algorithm compares and exchanges adjacent elements in the sequence to be sorted:

```
1.      procedure BUBBLE_SORT( $n$ )
2.      begin
3.          for  $i := n - 1$  downto 1 do
4.              for  $j := 1$  to  $i$  do
5.                  compare-exchange( $a_j, a_{j+1}$ );
6.      end BUBBLE_SORT
```

Sequential bubble sort algorithm.



Time Complexity


$$\text{Number of compare and exchange operations} = \sum_{i=1}^{n-1} i = \frac{n(n-1)}{2}$$

Indicates a time complexity of $O(n^2)$ given that a single compare-and-exchange operation has a constant complexity, $O(1)$.

Bubble Sort and its variants

- The complexity of bubble sort is $\Theta(n^2)$.
- Bubble sort is difficult to parallelize because the algorithm has no explicit parallelism!
- A simple variant (called Odd-Even Transposition), however, reveals an implicit parallelism ...

Odd-Even Transposition



Sorting $n = 8$ elements, using the odd-even transposition sort algorithm. During each phase, $n = 8$ elements are compared.

Odd-Even Transposition

```
1.  procedure ODD-EVEN( $n$ )
2.  begin
3.      for  $i := 1$  to  $n$  do
4.          begin
5.              if  $i$  is odd then
6.                  for  $j := 0$  to  $n/2 - 1$  do
7.                      compare-exchange( $a_{2j+1}, a_{2j+2}$ );
8.              if  $i$  is even then
9.                  for  $j := 1$  to  $n/2 - 1$  do
10.                     compare-exchange( $a_{2j}, a_{2j+1}$ );
11.          end for
12.  end ODD-EVEN
```

Sequential odd-even transposition sort algorithm.

Odd-Even Transposition

- After n phases of odd-even exchanges, the sequence is sorted.
- Each phase of the algorithm (either odd or even) requires $\Theta(n)$ comparisons.
- Serial complexity is $\Theta(n^2)$.

Parallel Odd-Even Transposition

- Consider the one item per processor case.
- There are n iterations, in each iteration, each processor does one compare-exchange.
- The parallel run time of this formulation is $\Theta(n)$.
- This is cost optimal with respect to the base serial algorithm but not the optimal one.

Parallel Odd-Even Transposition

```
1.  procedure ODD-EVEN_PAR( $n$ )
2.  begin
3.       $id :=$  process's label
4.      for  $i := 1$  to  $n$  do
5.          begin
6.              if  $i$  is odd then
7.                  if  $id$  is odd then
8.                      compare-exchange_min( $id + 1$ );
9.                  else
10.                     compare-exchange_max( $id - 1$ );
11.              if  $i$  is even then
12.                  if  $id$  is even then
13.                     compare-exchange_min( $id + 1$ );
14.                  else
15.                     compare-exchange_max( $id - 1$ );
16.          end for
17.  end ODD-EVEN_PAR
```

Parallel formulation of odd-even transposition.

Parallel Odd-Even Transposition

- Consider a block of n/p elements per processor.
- The first step is a local sort.
- In each subsequent step, the compare exchange operation is replaced by the compare split operation.
- The parallel run time of the formulation is

$$T_P = \overbrace{\Theta\left(\frac{n}{p} \log \frac{n}{p}\right)}^{\text{local sort}} + \overbrace{\Theta(n)}^{\text{comparisons}} + \overbrace{\Theta(n)}^{\text{communication}}.$$

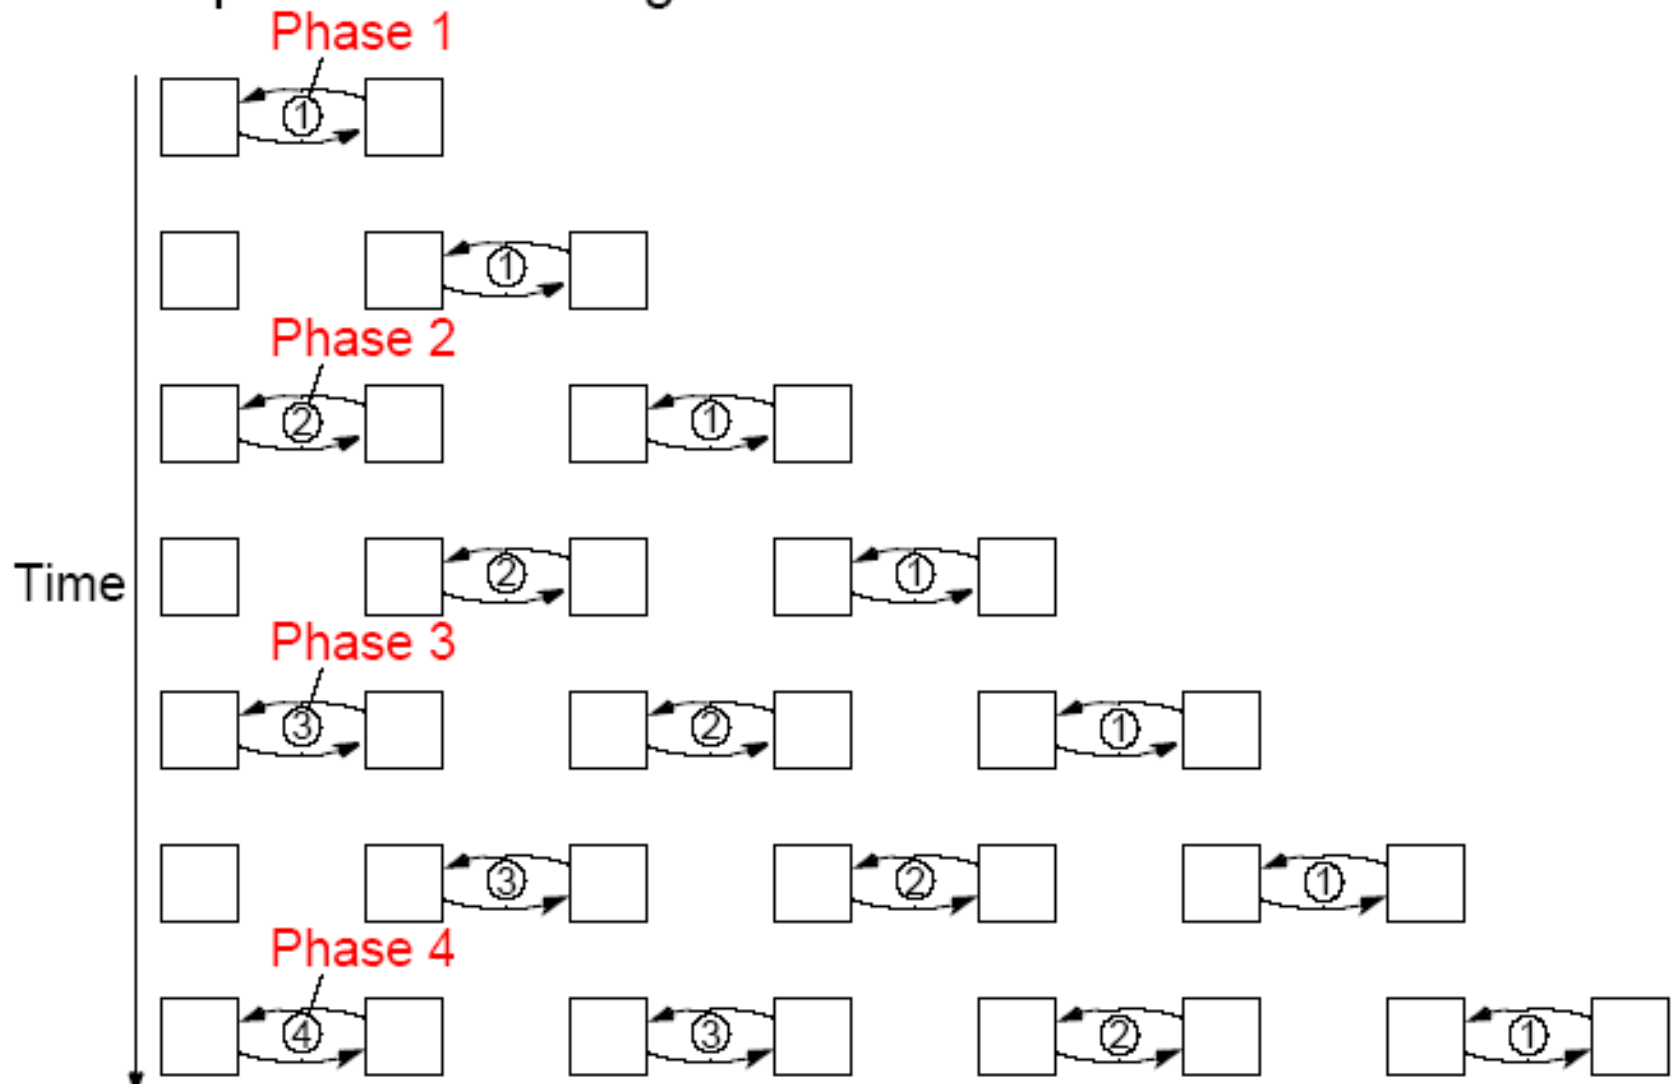
Parallel Odd-Even Transposition

- The parallel formulation is cost-optimal for $p = O(\log n)$.
- The isoefficiency function of this parallel formulation is $\Theta(p^2 2^p)$

(other slides on Parallel Odd-Even follow...)

Parallel Bubble Sort

Iteration could start before previous iteration finished if does not overtake previous bubbling action:



Odd-Even (Transposition) Sort

Variation of bubble sort.

Operates in two alternating phases, *even* phase and *odd* phase.

Even phase

Even-numbered processes exchange numbers with their right neighbor.

Odd phase

Odd-numbered processes exchange numbers with their right neighbor.

Odd-Even Transposition Sort

Sorting eight numbers

	P ₀	P ₁	P ₂	P ₃	P ₄	P ₅	P ₆	P ₇
Step								
0	4	↔ 2	7	↔ 8	5	↔ 1	3	↔ 6
1	2	↔ 4	7	↔ 8	1	↔ 5	3	6
2	2	↔ 4	7	↔ 1	8	↔ 3	5	↔ 6
3	2	↔ 4	1	↔ 7	3	↔ 8	5	6
4	2	↔ 1	4	↔ 3	7	↔ 5	8	↔ 6
5	1	↔ 2	3	↔ 4	5	↔ 7	6	8
6	1	↔ 2	3	↔ 4	5	↔ 6	7	↔ 8
7	1	↔ 2	3	↔ 4	5	↔ 6	7	8



Mergesort

A classical sequential sorting algorithm using divide-and-conquer approach. Unsorted list first divided into half. Each half is again divided into two. Continued until individual numbers are obtained.

Then pairs of numbers combined (merged) into sorted list of two numbers. Pairs of these lists of four numbers are merged into sorted lists of eight numbers. This is continued until the one fully sorted list is obtained.

Parallelizing Mergesort

Using tree allocation of processes



Analysis

Sequential

Sequential time complexity is $O(n \log n)$.

Parallel

$2 \log n$ steps in the parallel version but each step may need to perform more than one basic operation, depending upon the number of numbers being processed - see text.

Quicksort


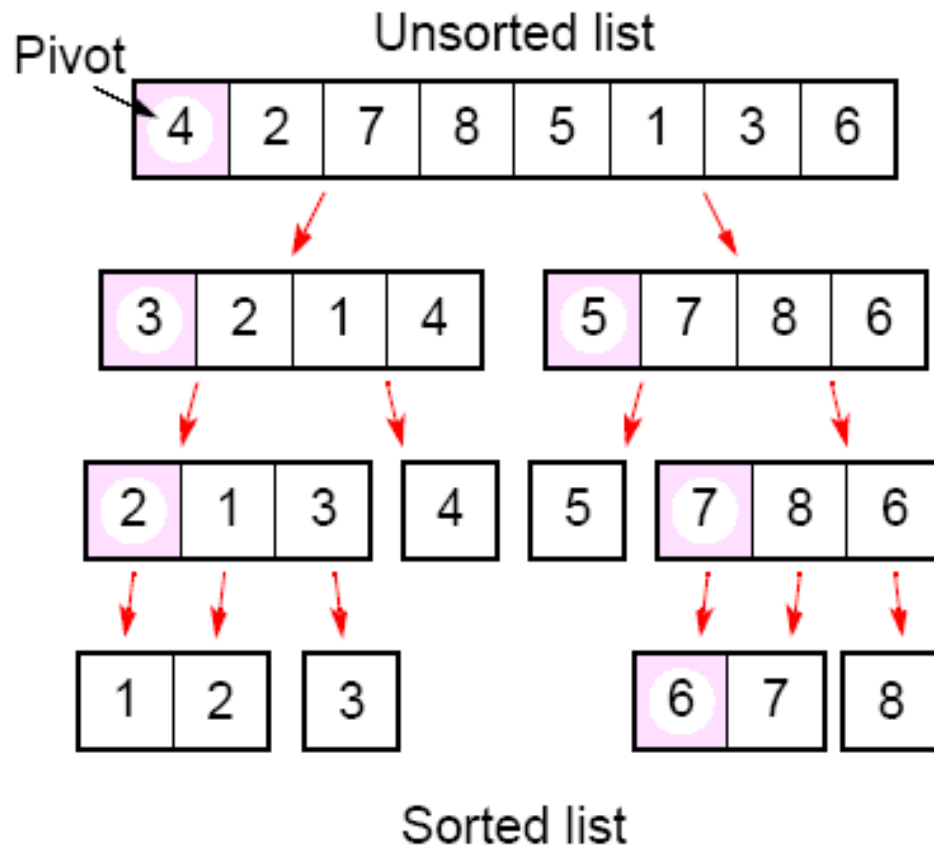
Very popular sequential sorting algorithm that performs well with an average sequential time complexity of $O(n \log n)$.

First list divided into two sublists. All the numbers in one sublist arranged to be smaller than all the numbers in the other sublist. Achieved by first selecting one number, called a *pivot*, against which every other number is compared. If the number is less than the pivot, it is placed in one sublist. Otherwise, it is placed in the other sublist.


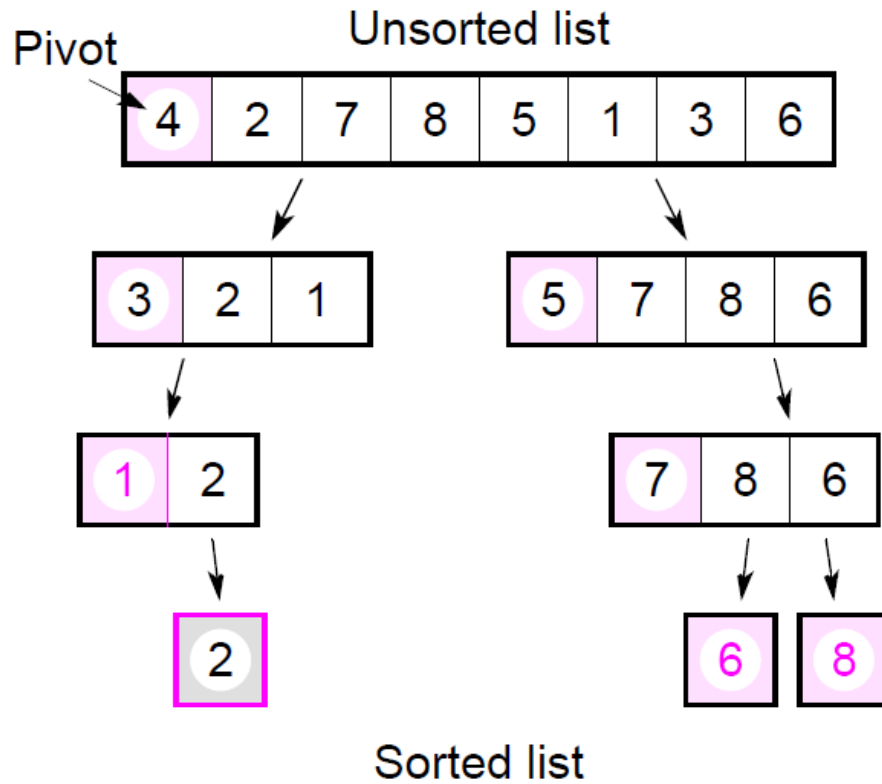
Pivot could be any number in the list, but often the first number in the list is chosen. Pivot itself could be placed in one sublist, or the pivot could be separated and placed in its final position.

Parallelizing Quicksort

Using tree allocation of processes



With the pivot being withheld in processes:




Analysis

Fundamental problem with all tree constructions – initial division done by a single processor, which will seriously limit speed.

Tree in quicksort will not, in general, be perfectly balanced Pivot selection very important to make quicksort operate fast.

Work Pool Implementation of Quicksort

First, work pool holds initial unsorted list. Given to first processor which divides list into two parts. One part returned to work pool to be given to another processor, while the other part operated upon again.



Neither Mergesort nor Quicksort parallelize very well as the processor efficiency is low (see book for analysis).

Quicksort also can be very unbalanced. Can use load balancing techniques.